

```
In [2]:
```

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LinearRegression
from sklearn.linear_model import LinearRegression, Ridge, Lasso
from sklearn.model_selection import GridSearchCV
from sklearn.metrics import r2_score, mean_squared_error

import warnings
warnings.filterwarnings('ignore')
```

```
In [3]: # Load the dataset
```

```
path = r"C:\Users\Downloads\archive (11)\uber.csv"
df = pd.read_csv(path)

# view dataset
df
```

```
Out[3]:
```

		Unnamed: 0	key	fare_amount	pickup_datetime	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff
0	24238194	2015-05-07 19:52:06.0000003		7.5	2015-05-07 19:52:06 UTC	-73.999817	40.738354	-73.999512	4
1	27835199	2009-07-17 20:04:56.0000002		7.7	2009-07-17 20:04:56 UTC	-73.994355	40.728225	-73.994710	4
2	44984355	2009-08-24 21:45:00.00000061		12.9	2009-08-24 21:45:00 UTC	-74.005043	40.740770	-73.962565	4
3	25894730	2009-06-26 08:22:21.0000001		5.3	2009-06-26 08:22:21 UTC	-73.976124	40.790844	-73.965316	4
4	17610152	2014-08-28 17:47:00.000000188		16.0	2014-08-28 17:47:00 UTC	-73.925023	40.744085	-73.973082	4
...
199995	42598914	2012-10-28 10:49:00.00000053		3.0	2012-10-28 10:49:00 UTC	-73.987042	40.739367	-73.986525	4
199996	16382965	2014-03-14 01:09:00.0000008		7.5	2014-03-14 01:09:00 UTC	-73.984722	40.736837	-74.006672	4
199997	27804658	2009-06-29 00:42:00.00000078		30.9	2009-06-29 00:42:00 UTC	-73.986017	40.756487	-73.858957	4
199998	20259894	2015-05-20 14:56:25.0000004		14.5	2015-05-20 14:56:25 UTC	-73.997124	40.725452	-73.983215	4
199999	11951496	2010-05-15 04:08:00.00000076		14.1	2010-05-15 04:08:00 UTC	-73.984395	40.720077	-73.985508	4

200000 rows × 9 columns

```
In [4]: df.columns
```

```
Out[4]: Index(['Unnamed: 0', 'key', 'fare_amount', 'pickup_datetime',
               'pickup_longitude', 'pickup_latitude', 'dropoff_longitude',
               'dropoff_latitude', 'passenger_count'],
              dtype='object')
```

```
In [5]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200000 entries, 0 to 199999
Data columns (total 9 columns):
 #   Column           Non-Null Count  Dtype  
 --- 
 0   Unnamed: 0        200000 non-null   int64  
 1   key              200000 non-null   object  
 2   fare_amount      200000 non-null   float64 
 3   pickup_datetime  200000 non-null   object  
 4   pickup_longitude 200000 non-null   float64 
 5   pickup_latitude  200000 non-null   float64 
 6   dropoff_longitude 199999 non-null   float64 
 7   dropoff_latitude  199999 non-null   float64 
 8   passenger_count  200000 non-null   int64  
dtypes: float64(5), int64(2), object(2)
memory usage: 13.7+ MB
```

```
In [6]: df.isnull().sum()
```

```
Out[6]: Unnamed: 0      0  
key          0  
fare_amount   0  
pickup_datetime 0  
pickup_longitude 0  
pickup_latitude 0  
dropoff_longitude 1  
dropoff_latitude 1  
passenger_count 0  
dtype: int64
```

```
In [7]: df = df.dropna()
```

```
In [8]: df.isnull().sum()
```

```
Out[8]: Unnamed: 0      0  
key          0  
fare_amount   0  
pickup_datetime 0  
pickup_longitude 0  
pickup_latitude 0  
dropoff_longitude 0  
dropoff_latitude 0  
passenger_count 0  
dtype: int64
```

```
In [9]: df['pickup_datetime'] = pd.to_datetime(df['pickup_datetime'])  
# print(df['pickup_datetime'])  
df['hour'] = df['pickup_datetime'].dt.hour  
# print(df['hour'])  
df['day_of_week'] = df['pickup_datetime'].dt.dayofweek  
# print(df['day_of_week'])
```

```
In [10]: # check datasets for more columns we added 'hour' and 'day_of_week' column  
df
```

```
Out[10]:
```

	Unnamed: 0	key	fare_amount	pickup_datetime	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	dropoff
0	24238194	2015-05-07 19:52:06.0000003	7.5	2015-05-07 19:52:06+00:00	-73.999817	40.738354	-73.999512	4	
1	27835199	2009-07-17 20:04:56.0000002	7.7	2009-07-17 20:04:56+00:00	-73.994355	40.728225	-73.994710	4	
2	44984355	2009-08-24 21:45:00.00000061	12.9	2009-08-24 21:45:00+00:00	-74.005043	40.740770	-73.962565	4	
3	25894730	2009-06-26 08:22:21.0000001	5.3	2009-06-26 08:22:21+00:00	-73.976124	40.790844	-73.965316	4	
4	17610152	2014-08-28 17:47:00.000000188	16.0	2014-08-28 17:47:00+00:00	-73.925023	40.744085	-73.973082	4	
...	
199995	42598914	2012-10-28 10:49:00.00000053	3.0	2012-10-28 10:49:00+00:00	-73.987042	40.739367	-73.986525	4	
199996	16382965	2014-03-14 01:09:00.0000008	7.5	2014-03-14 01:09:00+00:00	-73.984722	40.736837	-74.006672	4	
199997	27804658	2009-06-29 00:42:00.00000078	30.9	2009-06-29 00:42:00+00:00	-73.986017	40.756487	-73.858957	4	
199998	20259894	2015-05-20 14:56:25.0000004	14.5	2015-05-20 14:56:25+00:00	-73.997124	40.725452	-73.983215	4	
199999	11951496	2010-05-15 04:08:00.00000076	14.1	2010-05-15 04:08:00+00:00	-73.984395	40.720077	-73.985508	4	

199999 rows × 11 columns

```
In [11]: # Drop unnecessary columns  
df = df.drop(columns=['Unnamed: 0', 'key', 'pickup_datetime'])
```

```
In [12]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
Index: 199999 entries, 0 to 199999
Data columns (total 8 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   fare_amount    199999 non-null   float64
 1   pickup_longitude 199999 non-null   float64
 2   pickup_latitude 199999 non-null   float64
 3   dropoff_longitude 199999 non-null   float64
 4   dropoff_latitude 199999 non-null   float64
 5   passenger_count 199999 non-null   int64  
 6   hour          199999 non-null   int32  
 7   day_of_week    199999 non-null   int32  
dtypes: float64(5), int32(2), int64(1)
memory usage: 12.2 MB

```

```
In [13]: # check datasets for removal of columns we removed 'first_column with no name', 'key' and 'pickup_datetime' col
df
```

```
Out[13]:
```

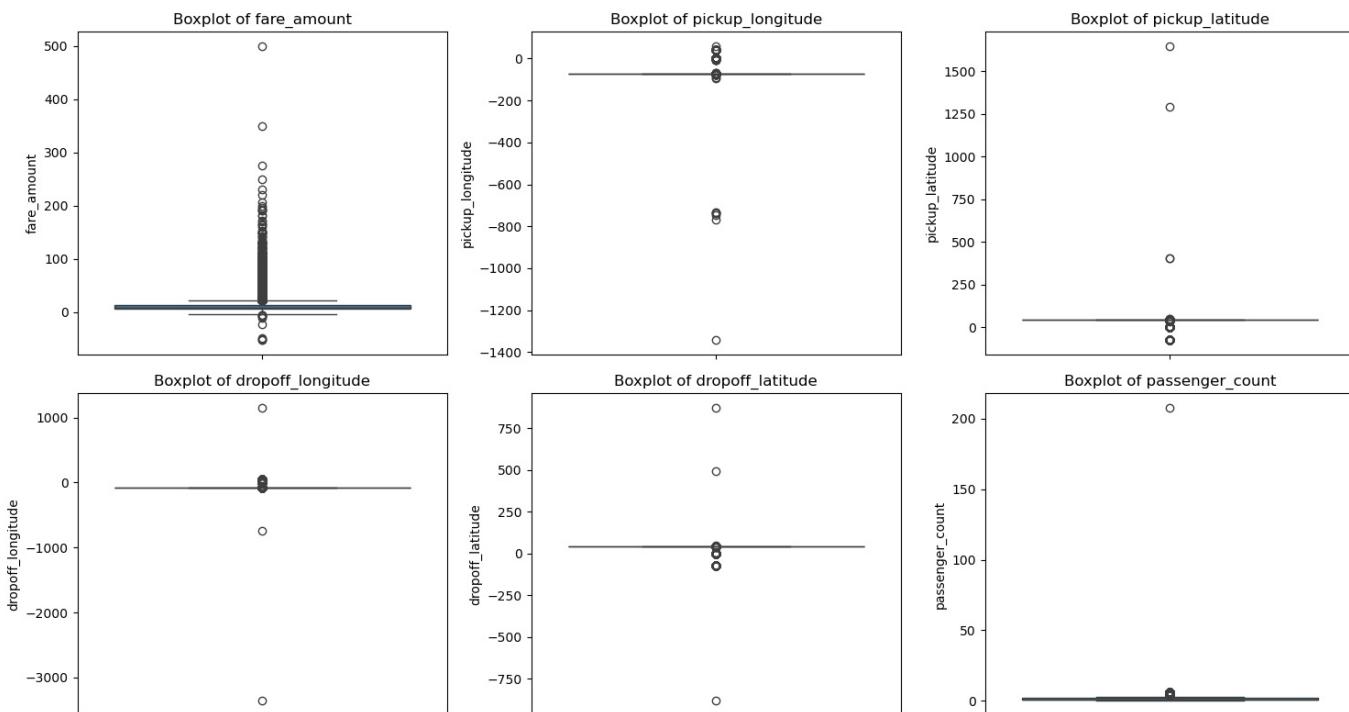
	fare_amount	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	passenger_count	hour	day_of_week
0	7.5	-73.999817	40.738354	-73.999512	40.723217		1	19
1	7.7	-73.994355	40.728225	-73.994710	40.750325		1	20
2	12.9	-74.005043	40.740770	-73.962565	40.772647		1	21
3	5.3	-73.976124	40.790844	-73.965316	40.803349		3	8
4	16.0	-73.925023	40.744085	-73.973082	40.761247		5	17
...
199995	3.0	-73.987042	40.739367	-73.986525	40.740297		1	10
199996	7.5	-73.984722	40.736837	-74.006672	40.739620		1	1
199997	30.9	-73.986017	40.756487	-73.858957	40.692588		2	0
199998	14.5	-73.997124	40.725452	-73.983215	40.695415		1	14
199999	14.1	-73.984395	40.720077	-73.985508	40.768793		1	4

199999 rows × 8 columns

```
In [14]: import matplotlib.pyplot as plt
import seaborn as sns

numeric_cols = ['fare_amount', 'pickup_longitude', 'pickup_latitude',
                'dropoff_longitude', 'dropoff_latitude', 'passenger_count']

plt.figure(figsize=(15, 8))
for i, col in enumerate(numeric_cols, 1):
    plt.subplot(2, 3, i)
    sns.boxplot(y=df[col])
    plt.title(f'Boxplot of {col}')
plt.tight_layout()
plt.show()
```



```
In [15]: # Split the data into features (X) and target (y)
X = df.drop(columns=['fare_amount']) # create new dataset ignoring 'fare_amount' column
y = df['fare_amount'] # create a series of only 'fare_amount' column
```

```
In [16]: # Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
In [17]: # Standardize the features (scaling)
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

```
In [18]: # Implement Linear Regression
lr_model = LinearRegression()
lr_model.fit(X_train_scaled, y_train)
y_pred_lr = lr_model.predict(X_test_scaled)
```

```
In [30]: # Ridge Regression
ridge = Ridge()
ridge_params = {'alpha': [0.01, 0.1, 1, 10, 100]}
ridge_grid = GridSearchCV(ridge, ridge_params, cv=5, scoring='r2')
ridge_grid.fit(X_train_scaled, y_train)
ridge_best = ridge_grid.best_estimator_
y_pred_ridge = ridge_best.predict(X_test_scaled)

# Lasso Regression
lasso = Lasso()
lasso_params = {'alpha': [0.01, 0.1, 1, 10, 100]}
lasso_grid = GridSearchCV(lasso, lasso_params, cv=5, scoring='r2')
lasso_grid.fit(X_train_scaled, y_train)
lasso_best = lasso_grid.best_estimator_
y_pred_lasso = lasso_best.predict(X_test_scaled)
```

```
In [38]: from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score

def evaluate_model(name, y_true, y_pred):
    r2 = r2_score(y_true, y_pred)
    rmse = np.sqrt(mean_squared_error(y_true, y_pred))
    mae = mean_absolute_error(y_true, y_pred)

    return {
        'Model': name,
        'R2 Score': r2,
        'RMSE': rmse,
        'MAE': mae
    }
```

```
In [40]: # Evaluate all models
results = []
results.append(evaluate_model('Linear Regression', y_test, y_pred_lr))
results.append(evaluate_model('Ridge Regression', y_test, y_pred_ridge))
results.append(evaluate_model('Lasso Regression', y_test, y_pred_lasso))

results_df = pd.DataFrame(results)
print("\nModel Evaluation Results:")
print(results_df)
```

```
Model Evaluation Results:
      Model   R2 Score       RMSE       MAE
0  Linear Regression  0.000914  10.194550  6.064055
1  Ridge Regression  0.000914  10.194551  6.064055
2  Lasso Regression  0.000870  10.194776  6.064411
```

```
In [ ]:
```